



# **MALIGNANT COMMENTS CLASSIFICATION PROJECT**

Submitted by  
NARRALA HOMAKIRAN

## **ACKNOWLEDGMENT**

The project entitled “MALIGNANT COMMENTS CLASSIFICATION” is done by me during my internship with Flip Robo Technologies. I am grateful to Data Trained and Flip Robo Technologies for their guidance during this project.

Other reference websites used to complete this project are:

1. Data source provided by the client.
2. Wikipedia
3. Towardsdatascience.com
4. Datacamp.com

# **INTRODUCTION**

## **Conceptual Background of the Domain Problem**

The proliferation of social media enables people to express their opinions widely online. However, at the same time, this has resulted in the emergence of conflict and hate, making online environments uninviting for users. Although researchers have found that hate is a problem across multiple platforms, there is a lack of models for online hate detection.

Online hate, described as abusive language, aggression, cyberbullying, hatefulness and many others has been identified as a major threat on online social media platforms. Social media platforms are the most prominent grounds for such toxic behaviour.

There has been a remarkable increase in the cases of cyberbullying and trolls on various social media platforms. Many celebrities and influences are facing backlashes from people and have to come across hateful and offensive comments. This can take a toll on anyone and affect them mentally leading to depression, mental illness, self-hatred and suicidal thoughts.

Internet comments are bastions of hatred and vitriol. While online anonymity has provided a new outlet for aggression and hate speech, machine learning can be used to fight it. The problem we sought to solve was the tagging of internet comments that are aggressive towards other users. This means that insults to third parties such as celebrities will be tagged as unoffensive, but “u are an idiot” is clearly offensive.

## **PROBLEM STATEMENT:**

Our goal is to build a prototype of online hate and abuse comment classifier which can be used to classify hate and offensive comments so that it can be controlled and restricted from spreading hatred and cyberbullying.

# Analytical Problem Framing

## Mathematical/ Analytical Modeling of the Problem

- With continuous increase in available data, there is a pressing need to organize it and modern classification problems often involve the prediction of multiple labels simultaneously associated with a single instance. Known as Multi-Label Classification, it is one such task which is omnipresent in many real world problems. In this project also, we have multi-label classification problem.
- We have used Tf-Idf Vectorizer to vectorize the words in our dataset. TF-IDF is an abbreviation for Term Frequency Inverse Document Frequency. This is very common algorithm to transform text into a meaningful representation of numbers which is used to fit machine algorithm for prediction. It is very important for tuning performance on NLP projects.
- The TF-IDF score for the word  $t$  in the document  $d$  from the document set  $D$  is calculated as follows:

$$tfidf(t, d, D) = tf(t, d) \cdot idf(t, D)$$

- Where:

$$tf(t, d) = \log(1 + freq(t, d))$$

$$idf(t, D) = \log\left(\frac{N}{count(d \in D: t \in d)}\right)$$

## Data Sources and their formats

- The data set contains the training set, which has approximately 1,59,000 samples and the test set which contains nearly 1,53,000 samples. All the data samples contain 8 fields which includes 'Id', 'Comments', 'Malignant', 'Highly malignant', 'Rude', 'Threat', 'Abuse' and 'Loathe'.
- The label can be either 0 or 1, where 0 denotes a NO while 1 denotes a YES. There are various comments which have multiple labels. The first attribute is a unique ID associated with each comment.

The data set includes:

- **Malignant:** It is the Label column, which includes values 0 and 1, denoting if the comment is malignant or not.
- **Highly Malignant:** It denotes comments that are highly malignant and hurtful.
- **Rude:** It denotes comments that are very rude and offensive.
- **Threat:** It contains indication of the comments that are giving any threat to someone.
- **Abuse:** It is for comments that are abusive in nature.
- **Loathe:** It describes the comments which are hateful and loathing in nature.
- **ID:** It includes unique Ids associated with each comment text given.
- **Comment text:** This column contains the comments extracted from various social media platforms.

The sample data for the reference is as shown below:

```
Out[4]:
```

	id	comment_text	malignant	highly_malignant	rude	threat	abuse	loathe
0	0000907932d777bf	ExplanationWhy the edits made under my usern...	0	0	0	0	0	0
1	0001030d9cb6b0f	Draw! He matches this background colour I'm s...	0	0	0	0	0	0
2	00011307e0026f	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0
3	0001b41b1c0bb37e	"inMore! can't make any real suggestions on ...	0	0	0	0	0	0
4	0001d958c540b635	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0
5	0002548504725e87	"inCongratulations from me as well, use the ...	0	0	0	0	0	0
6	0002b2b3a9d6b337	COCKSUCKER BEFORE YOU PISS AROUND ON MY WORK	1	1	1	0	1	0
7	00031b1e95af7021	Your vandalism to the Matt Shrivington article...	0	0	0	0	0	0
8	000372b1f536c51d	Sorry if the word 'nonsense' was offensive to ...	0	0	0	0	0	0
9	0004093b2687caa	alignment on this subject and which are contra...	0	0	0	0	0	0

```
In [5]: df_test.head(10)
```

```
Out[5]:
```

	id	comment_text
0	000010ee341f8b12	Ye bish Ja Rule is more successful than you'll...
1	0000247987823ef7	== From RBC == in The title is fine as it is...
2	00013b17ad223c48	" in Sources == in " Zane Ashton on Lap...
3	00017563c977919a	:If you have a look back at the source, the in...
4	00017969a6997eb	I don't anonymously edit articles at all.
5	0001ea871776ae05	Thank you for understanding. I think very high...
6	0002411554db0e0f	Please do not add nonsense to Wikipedia. Such ...
7	000247e63dcd1211	:Dear god this site is horrible.
8	00025356b4737918	" In Only a fool can believe in such numbers. ...
9	00026b1002e71cc	== Double Redirects == in When fixing double...

- Then we further checked more about data using info, shapes using .shape, columns using .columns(), null values using .isnull().sum(), and further visualize it through heatmap as follows:

```

In [6]: df_train.shape
Out[6]: (159571, 8)

In [7]: df_test.shape
Out[7]: (153164, 2)

In [8]: df_train.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 159571 entries, 0 to 159570
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                     159571 non-null object
1   comment_text           159571 non-null object
2   malignant               159571 non-null int64
3   highly_malignant        159571 non-null int64
4   rude                   159571 non-null int64
5   threat                 159571 non-null int64
6   abuse                  159571 non-null int64
7   loathe                 159571 non-null int64
dtypes: int64(6), object(2)
memory usage: 9.7+ MB

In [9]: df_train.columns
Out[9]: Index(['id', 'comment_text', 'malignant', 'highly_malignant', 'rude', 'threat',
              'abuse', 'loathe'],
              dtype='object')

In [10]: df_train.isnull().sum()
Out[10]: id                0
comment_text              0
malignant                 0
highly_malignant         0
rude                     0
threat                   0
abuse                    0
loathe                   0
dtype: int64

In [11]: sns.heatmap(df_train.isnull())
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x2d84d041738>

```



- Then we perform Exploratory Data Analysis(EDA) as follows:

```

In [12]: # Dropping 'id' to reduce unnecessary feature
df_train.drop(['id'],axis=1,inplace=True)


In [13]: df_train.describe()
Out[13]:
          malignant  highly_malignant      rude  threat      abuse  loathe
count  159571.000000    159571.000000  159571.000000  159571.000000  159571.000000  159571.000000
mean      0.085844      0.009990      0.052948      0.002996      0.049364      0.008805
std       0.294379      0.090477      0.223931      0.054850      0.219827      0.093420
min       0.000000      0.000000      0.000000      0.000000      0.000000      0.000000
25%      0.000000      0.000000      0.000000      0.000000      0.000000      0.000000
50%      0.000000      0.000000      0.000000      0.000000      0.000000      0.000000
75%      0.000000      0.000000      0.000000      0.000000      0.000000      0.000000
max       1.000000      1.000000      1.000000      1.000000      1.000000      1.000000





In [14]: features=df_train.columns[1:]
features
Out[14]: Index(['malignant', 'highly_malignant', 'rude', 'threat', 'abuse', 'loathe'], dtype='object')

In [15]: for i in features:
print('Number of unique values in {}: {}'.format(i, df_train[i].nunique()))

Number of unique values in malignant : 2
Number of unique values in highly_malignant : 2
Number of unique values in rude : 2
Number of unique values in threat : 2
Number of unique values in abuse : 2
Number of unique values in loathe : 2

```





```
In [17]: for i in features:
          print('Number of unique values in {}: {}'.format(i, df_train[i].value_counts()))

Number of unique values in malignant : 0    144277
                                         1    15294
Name: malignant, dtype: int64
Number of unique values in highly_malignant : 0    157976
                                                1    1595
Name: highly_malignant, dtype: int64
Number of unique values in rude : 0    151122
                                   1     8449
Name: rude, dtype: int64
Number of unique values in threat : 0    159093
                                    1      478
Name: threat, dtype: int64
Number of unique values in abuse : 0    151694
                                    1     7877
Name: abuse, dtype: int64
Number of unique values in loathe : 0    158166
                                    1     1405
Name: loathe, dtype: int64
```

```
In [18]: good_comments = df_train[(df_train['malignant']!=1) & (df_train['highly_malignant']!=1) & (df_train['rude']!=1) &
          (df_train['threat']!=1) & (df_train['abuse']!=1) & (df_train['loathe']!=1)]
good_percent=len(good_comments)/len(df_train)*100
print('Percentage of good comments = ',good_percent)
print('Percentage of negative comments = ', (100-good_percent))

Percentage of good comments = 89.83211235124176
Percentage of negative comments = 10.167887648758239
```

```
In [19]: df_train['comment_length']=df_train.comment_text.str.len()
df_train
```

```
Out[19]:
```

	comment_text	malignant	highly_malignant	rude	threat	abuse	loathe	comment_length
0	ExplanationWhy the edits made under my usern...	0	0	0	0	0	0	264
1	D'aww! He matches this background colour I'm s...	0	0	0	0	0	0	112
2	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0	233
3	"nMore!nI can't make any real suggestions on ...	0	0	0	0	0	0	622
4	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0	67
...	...	...	...	...	...	...	...	...
159566	"....And for the second time of asking, when ...	0	0	0	0	0	0	295
159567	You should be ashamed of yourself in/nThat is ...	0	0	0	0	0	0	99
159568	Spitzer in/nUmm, theres no actual article for ...	0	0	0	0	0	0	81
159569	And it looks like it was actually you who put ...	0	0	0	0	0	0	116
159570	"nAnd ... I really don't think you understand...	0	0	0	0	0	0	189

159571 rows x 8 columns

```
In [20]: df_test['comment_length']=df_test.comment_text.str.len()
df_test.head()
```

```
Out[20]:
```

	id	comment_text	comment_length
0	000010ee341fdb12	Yo bltch Ja Rule is more succesful then you'll...	387
1	0000247807823e7	== From RIC == in/n The tile is fine as it is...	50
2	00013b17ad220e46	" in/n == Sources == in/n * Zawe Ashton on Lap...	54
3	00017563c37f919a	:If you have a look back at the source, the in...	205
4	00017895a0897eb	I don't anonymously edit articles at all.	41

```
In [21]: # Creating a new feature having Negative Comments and Non-Negative Comments.
df_train['label'] = df_train[features].max(axis=1)
df_train.head(10)
```

```
Out[21]:
```

	comment_text	malignant	highly_malignant	rude	threat	abuse	loathe	comment_length	label
0	ExplanationWhy the edits made under my usern...	0	0	0	0	0	0	264	0
1	D'aww! He matches this background colour I'm s...	0	0	0	0	0	0	112	0
2	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0	233	0
3	"nMore!nI can't make any real suggestions on ...	0	0	0	0	0	0	622	0
4	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0	67	0
5	"n/nCongratulatlons from me as well, use the ...	0	0	0	0	0	0	65	0
6	COCKSUCKER BEFORE YOU PISS AROUND ON MY WORK	1	1	1	0	1	0	44	1
7	Your vandallism to the Matt Shirvington article...	0	0	0	0	0	0	115	0
8	Sorry if the word 'nonsense' was offensive to ...	0	0	0	0	0	0	472	0
9	alignment on this subject and which are contra...	0	0	0	0	0	0	70	0

```
In [22]: df_train['label'].value_counts()
```

```
Out[22]: 0    143346
         1     16225
         Name: label, dtype: int64
```

- Next we will be using a function to filter using POS tagging. Also, all the pre-processing steps needed to clean the data.

```
In [24]: def get_pos(pos_tag):
if pos_tag.startswith('J'):
    return wordnet.ADJ
elif pos_tag.startswith('N'):
    return wordnet.NOUN
elif pos_tag.startswith('R'):
    return wordnet.ADV
else:
    return wordnet.NOUN
```

```
In [25]: # Function for data cleaning...
def Processed_data(comments):
    # Replace email addresses with 'email'
    comments=re.sub(r'[\w\d\.-]+\@[\w\d\.-]+\.[a-z]{2,}$', ' ', comments)

    # Replace 10 digit phone numbers (formats include paranthesis, spaces, no spaces, dashes) with 'phonenumber'
    comments=re.sub(r'(\d{3})?(\d{3})?(\d{3})?(\d{3})?(\d{3})?(\d{3})?(\d{3})?(\d{3})?(\d{3})?(\d{3})?$', ' ', comments)

    # getting only words(i.e removing all the special characters)
    comments = re.sub(r'[\W_]', ' ', comments)

    # getting only words(i.e removing all the " _ ")
    comments = re.sub(r'[_ ]', ' ', comments)

    # getting rid of unwanted characters(i.e remove all the single characters left)
    comments=re.sub(r'[\s+[\a-zA-Z]\s+]', ' ', comments)

    # Removing extra whitespaces
    comments=re.sub(r'\s+', ' ', comments, flags=re.I)

    #converting all the letters of the review into lowercase
    comments = comments.lower()

    # splitting every words from the sentences
    comments = comments.split()

    # iterating through each words and checking if they are stopwords or not,
    comments=[word for word in comments if not word in set(STOPWORDS)]

    # remove empty tokens
    comments = [text for text in comments if len(text) > 0]

    # getting pos tag text
    pos_tags = pos_tag(comments)

    # considering words having length more than 3only
    comments = [text for text in comments if len(text) > 3]

    # performing Lemmatization operation and passing the word in get_pos function to get filtered using POS ...
    comments = [(WordNetLemmatizer().lemmatize(text[i]), get_pos(text[i]))for text in pos_tags]

    # considering words having length more than 3 only
    comments = [text for text in comments if len(text) > 3]
    comments = ' '.join(comments)
    return comments
```

- After performing all the above steps and also adding a new feature to check new comment length after cleaning, our dataset would look as follows:



```
In [28]: # Adding new feature clean_comment_length to store length of characters
df_train['clean_comment_length'] = df_train['clean_comment_text'].apply(lambda x: len(str(x)))
df_train.head()
```

```
Out[28]:
```


	comment_text	malignant	highly_malignant	rude	threat	abuse	loathe	comment_length	label	clean_comment_text	clean_comment_length
0	Explanation/Why the edits made under my usern...	0	0	0	0	0	0	264	0	explanation edits username hardcore metallica ...	129
1	'Draww! He matches this background colour I'm s...	0	0	0	0	0	0	112	0	match background colour seemingly stuck tracks...	64
2	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0	233	0	trying edit constantly removing relevant infor...	112
3	"InMoreInI can't make any real suggestions on ...	0	0	0	0	0	0	622	0	real suggestion improvement wondered section a...	315
4	'You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0	67	0	hero chance remember page	26

```
In [29]: # Adding new feature clean_comment_length to store length of characters
df_test['clean_comment_length'] = df_test['clean_comment_text'].apply(lambda x: len(str(x)))
df_test.head()
```

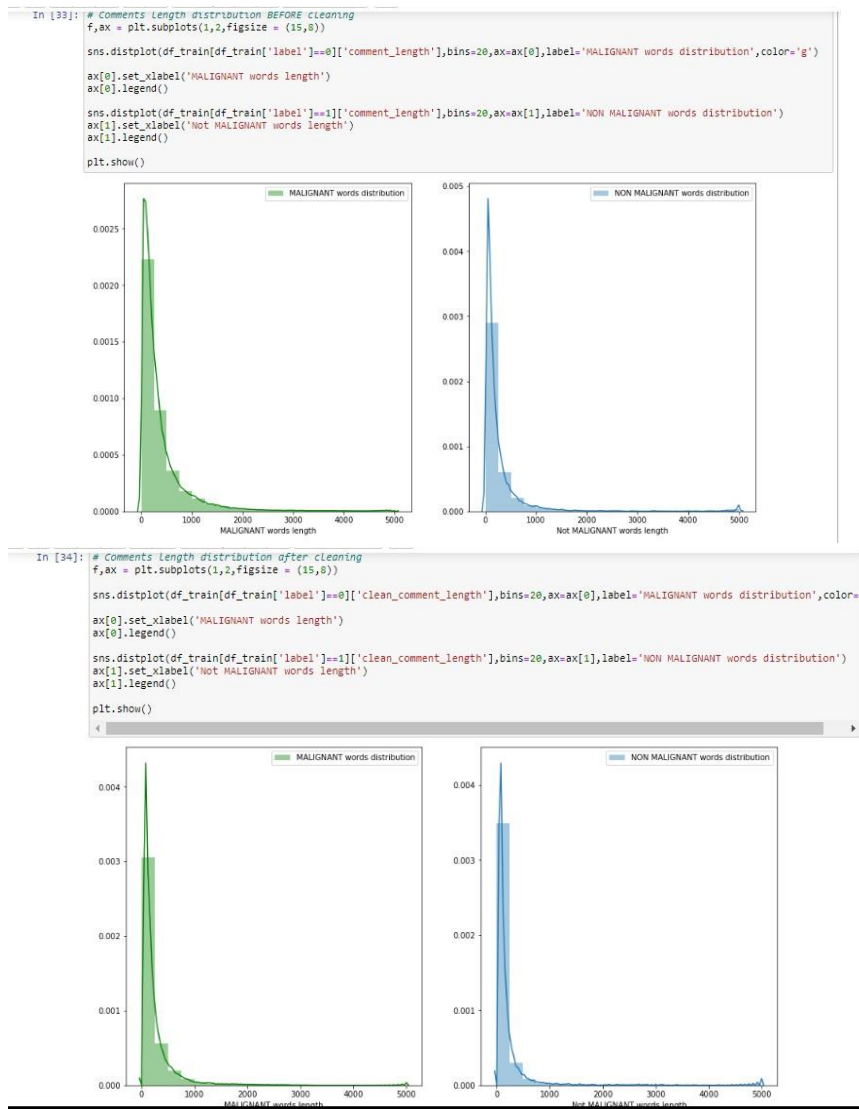
```
Out[29]:
```

	id	comment_text	comment_length	clean_comment_text	clean_comment_length
0	000010ee341fdb12	Yo bitch Ja Rule is more succesful then you'll...	367	bitch rule succesful whats hating mofuckas bit...	184
1	0000247867823ef7	== From RRC == 'n/n The title is fine as it is...	50	title fine	10
2	00013b17ad220c46	" 'n/n == Sources == 'n/n * Zawe Ashton on Lap...	54	source zawe ashton lapland	26
3	00017563c37f919a	:if you have a look back at the source, the in...	205	look source information updated correct form g...	109
4	00017565ad8997eb	I don't anonymously edit articles at all.	41	anonymously edit article	24

- We have also observed most frequent words in positive and negative comments through word-cloud:



- Then we have checked distribution of comment length before and after cleaning.



## Preparing Data For Modelling

- We are using TF-IDF vectorizer for vectorizing the words.

```

In [35]: # TF-IDF(term frequency-inverse document frequency) vectorizer
def Tf_idf_train(text):
    tfidf = TfidfVectorizer(min_df=3,smooth_idf=False)
    return tfidf.fit_transform(text)

```

```

In [36]: # Let's define x, y for modelling
x=Tf_idf_train(df_train['clean_comment_text'])
x.shape

```

Out[36]: (159571, 43194)

```

In [37]: # For y
y = df_train['label'].values
y.shape

```

Out[37]: (159571,)

# MODEL BUILDING

```
In [38]: # Importing useful Libraries for model training

from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier

# Model selection libraries...
from sklearn.model_selection import cross_val_score, cross_val_predict, train_test_split
from sklearn.model_selection import GridSearchCV

# Importing some metrics we can use to evaluate our model performance....
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.metrics import roc_auc_score, roc_curve, auc
from sklearn.metrics import precision_score, recall_score, f1_score

In [39]: #splitting the data into training and testing
X_train,X_test,y_train,y_test=train_test_split(x,y,random_state=42,test_size=0.30,stratify=y)

In [40]: # Creating instances for different Classifiers

LR=LogisticRegression()
MNB=MultinomialNB()
DT=DecisionTreeClassifier()
KNN=KNeighborsClassifier()
SV=SVC()

In [41]: # Putting Scikit-Learn machine Learning Models in a list so that it can be used for further evaluation in loop.
models=[]
models.append(('LogisticRegression',LR))
models.append(('MultinomialNB',MNB))
models.append(('DecisionTreeClassifier',DT))
models.append(('KNeighborsClassifier',KNN))
models.append(('SVC',SV))
```

- We obtained following results after training the model on various algorithms:

```
***** LogisticRegression *****

LogisticRegression()

Accuracy_score= 0.9531667780748663

Cross_Val_Score= 0.9535128562209592

roc_auc_score= 0.7924918146002661

classification report
precision    recall  f1-score   support

   0       0.96       0.99       0.97       43804
   1       0.92       0.59       0.72        4868

 accuracy         0.94
 macro avg       0.94       0.79       0.85
weighted avg       0.95       0.95       0.95       47672

[[42754  250]
 [ 1992 2876]]

AxesSubplot(0.125,0.808774;0.62x0.0712264)
```

```

***** MultinomialNB *****

MultinomialNB()

Accuracy_score= 0.9354737633689839

Cross_Val_Score= 0.9367383554748633

roc_auc_score= 0.6884622511658735

classification report
precision    recall  f1-score   support

     0       0.93      1.00      0.97    43004
     1       0.97      0.38      0.54     4868

 accuracy
macro avg       0.95      0.69      0.75    47872
weighted avg     0.94      0.94      0.92    47872

[[42941  63]
 [ 3826 1842]]

AxesSubplot(0.125,0.808774;0.62x0.0712264)

```

```

***** DecisionTreeClassifier *****

DecisionTreeClassifier()

Accuracy_score= 0.9392337901069518

Cross_Val_Score= 0.9399076251956352

roc_auc_score= 0.8263624575788062

classification report
precision    recall  f1-score   support

     0       0.96      0.97      0.97    43004
     1       0.71      0.68      0.70     4868

 accuracy
macro avg       0.84      0.83      0.83    47872
weighted avg     0.94      0.94      0.94    47872

[[41630 1374]
 [ 1535 3333]]

AxesSubplot(0.125,0.808774;0.62x0.0712264)

```

```

***** KNeighborsClassifier *****

KNeighborsClassifier()

Accuracy_score= 0.8963277072192514

Cross_Val_Score= 0.8967356214680471

roc_auc_score= 0.6214955391587276

classification report
precision    recall  f1-score   support

     0       0.92      0.97      0.94    43004
     1       0.48      0.28      0.35     4868

 accuracy
macro avg       0.70      0.62      0.65    47872
weighted avg     0.88      0.90      0.88    47872

[[41563 1441]
 [ 3522 1346]]

AxesSubplot(0.125,0.808774;0.62x0.0712264)

```

```

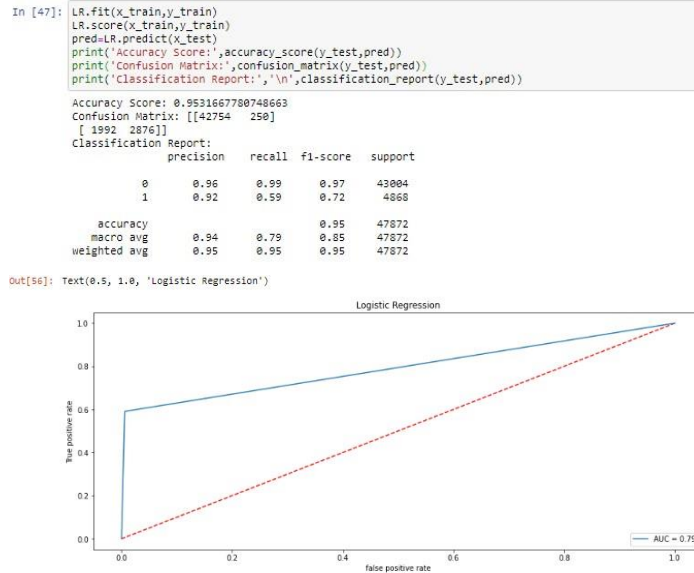
***** SVC *****

SVC()

Accuracy_score= 0.9545872326203209

```

- Since the dataset was too large it took me around 9-10 hours to get these results for these algorithms. Hence I had to interrupt the kernel and proceed with available results. Out of all the algorithms, Logistic Regression was giving best score. Also, its cross validation score was also satisfactory. Its ROC\_AUC curve is as shown:



## PREDICTING TEST DATASET

```
In [51]: def Tf_idf_test(text):
tfidf = TfidfVectorizer(max_features=43194,smooth_idf=False)
return tfidf.fit_transform(text)
```

```
In [52]: x_test_data=Tf_idf_test(df_test['clean_comment_text'])
```

```
In [53]: x_test_data.shape
```

Out[53]: (153164, 43194)

```
In [54]: Prediction=LR.predict(x_test_data)
df_test['Predicted Labels']=Prediction
df_test
```

Out[54]:

	id	comment_text	comment_length	clean_comment_text	clean_comment_length	Predicted Labels
0	000010ee341fdb12	Yo bitch Ja Rule is more succesful than you'll...	367	bitch rule succesful whats hating mofuckas bit...	184	0
1	0000247867823ef7	== From RfC == \n\n The title is fine as it is...	50	title fine	10	0
2	00013b17ad220c48	" \n\n == Sources == \n\n " Zawe Ashton on Lap...	54	source zawe ashton lapland	26	0
3	00017563c37f919a	:if you have a look back at the source, the in...	205	look source information updated correct form g...	109	0
4	00017685ad8997eb	I don't anonymously edit articles at all.	41	anonymously edit article	24	0
...	...	...	...	...	...	...
153159	ffcd0900ee309b5	. \n i totally agree, this stuff is nothing bu...	60	totally agree stuff long crap	29	0
153160	ffcd7a9a9eb32c16	== Throw from out field to home plate. == \n\n...	198	throw field home plate faster throwing direct ...	85	0
153161	ffda9e8d9afa9e	" \n\n == Okinotorishima categories == \n\n   ...	423	okinotorishima category change agree correct g...	212	0

```
In [48]: df_test['Predicted Labels'].value_counts()
```

Out[48]: 0 152452  
1 712  
Name: Predicted Labels, dtype: int64

```
In [49]: # Pickle file.
import joblib
joblib.dump(LR, 'Malignant_Prediction.pkl')
```

Out[49]: ['Malignant\_Prediction.pkl']

```
In [50]: df_test.to_csv('Malignant_Predict.csv')
```

## **CONCLUSION**

- We have got Logistic Regression as best model since it's giving us good result and other metrics are also satisfactory.
- Using Logistic Regression as our final algorithm we have predicted the values for test dataset and it's also working well and is able to differentiate/predict negative comments and non-negative (good) comments.